

Loan Application Status Prediction

Problem Statement:

This dataset includes details of applicants who have applied for loan. The dataset includes details like credit history, loan amount, their income, dependents etc.

Independent Variables:

- Loan_ID
- Gender
- Married
- Dependents
- Education
- Self_Employed
- ApplicantIncome
- CoapplicantIncome
- Loan_Amount
- Loan_Amount_Term
- Credit History
- Property_Area

Dependent Variable (Target Variable):

- Loan_Status

```
In [ ]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn import svm
from sklearn.metrics import accuracy_score
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: loan_df=pd.read_csv('Loan_Application_Status_Prediction.csv')
loan_df.head(10)
```

```
Out[2]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplicant
0	LP001002	Male	No	0	Graduate	No	5849	
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	
5	LP001011	Male	Yes	2	Graduate	Yes	5417	
6	LP001013	Male	Yes	0	Not Graduate	No	2333	
7	LP001014	Male	Yes	3+	Graduate	No	3036	
8	LP001018	Male	Yes	2	Graduate	No	4006	
9	LP001020	Male	Yes	1	Graduate	No	12841	

```
In [3]: #number of rows and columns
loan_df.shape
```

```
Out[3]: (614, 13)
```

```
In [4]: #statistical measures
loan_df.describe()
```

```
Out[4]:
```

	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	614.000000	614.000000	592.000000	600.000000	564.000000
mean	5403.459283	1621.245798	146.412162	342.000000	0.842199
std	6109.041673	2926.248369	85.587325	65.12041	0.364878
min	150.000000	0.000000	9.000000	12.000000	0.000000
25%	2877.500000	0.000000	100.000000	360.000000	1.000000
50%	3812.500000	1188.500000	128.000000	360.000000	1.000000
75%	5795.000000	2297.250000	168.000000	360.000000	1.000000
max	81000.000000	41667.000000	700.000000	480.000000	1.000000

```
In [5]: #number of missing values in each columns
loan_df.isnull().sum()
```

```
Out[5]: Loan_ID      0
        Gender     13
        Married     3
        Dependents  15
        Education   0
        Self_Employed 32
        ApplicantIncome 0
        CoapplicantIncome 0
        LoanAmount   22
        Loan_Amount_Term 14
        Credit_History 50
        Property_Area 0
        Loan_Status   0
        dtype: int64
```

```
In [6]: #dropping the missing values
        loan_df=loan_df.dropna()
```

```
In [7]: #Again check number of missing values in each column
        loan_df.isnull().sum()
```

```
Out[7]: Loan_ID      0
        Gender      0
        Married      0
        Dependents   0
        Education    0
        Self_Employed 0
        ApplicantIncome 0
        CoapplicantIncome 0
        LoanAmount   0
        Loan_Amount_Term 0
        Credit_History 0
        Property_Area 0
        Loan_Status   0
        dtype: int64
```

Here no missing values found

```
In [8]: # Label encoding
        loan_df.replace({"Loan_Status":{"N":0,'Y':1}},inplace=True)
```

```
In [9]: loan_df.head()
```


```
Out[9]:
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplicant
1	LP001003	Male	Yes	1	Graduate	No	4583	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	
4	LP001008	Male	No	0	Graduate	No	6000	
5	LP001011	Male	Yes	2	Graduate	Yes	5417	

Correlation

```
In [33]: plt.figure(figsize=(15,12))
sns.heatmap(loan_df.corr(),cbar=True,square=True,fmt='.2f',annot=True,annot_kws={'size
```

```
Out[33]: <AxesSubplot:>
```



```
In [10]: #Dependent columns values
loan_df['Dependents'].value_counts()
```

```
Out[10]: 0    274
         2     85
         1     80
         3+    41
         Name: Dependents, dtype: int64
```

```
In [11]: #replacing the value of 3+ to 4
loan_df=loan_df.replace(to_replace='3+',value=4)
```


```
In [12]: #Dependent values
loan_df['Dependents'].value_counts()
```

```
Out[12]: 0    274
         2     85
         1     80
         4     41
         Name: Dependents, dtype: int64
```

Data Visualization


```
In [13]: # Education & Loan Status  
sns.countplot(x='Education', hue='Loan_Status', data=loan_df)
```

```
Out[13]: <AxesSubplot:xlabel='Education', ylabel='count'>
```



```
In [14]: #Marital status & Loan Status  
sns.countplot(x='Married', hue='Loan_Status', data=loan_df)
```

```
Out[14]: <AxesSubplot:xlabel='Married', ylabel='count'>
```



```
In [15]: # convert categorical columns to numerical values  
loan_df.replace({'Married':{'No':0, 'Yes':1}, 'Gender':{'Male':1, 'Female':0}, 'Self_Employed':{'No':0, 'Yes':1},  
                'Property_Area':{'Rural':0, 'Semiurban':1, 'Urban':2}, 'Education': {'Graduate':1, 'Not Graduate':0}})
```

```
In [16]: loan_df.head()
```

Out[16]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coapplicant
1	LP001003	1	1	1	1	0	4583	
2	LP001005	1	1	0	1	1	3000	
3	LP001006	1	1	0	0	0	2583	
4	LP001008	1	0	0	1	0	6000	
5	LP001011	1	1	2	1	1	5417	

```
In [17]: # separating the data and label
x = loan_df.drop(columns=['Loan_ID','Loan_Status'],axis=1)
y = loan_df['Loan_Status']
```

```
In [18]: print(x.shape)
print(y.shape)
```

```
(480, 11)
```

```
(480,)
```

Train Test Split

```
In [19]: x_train, x_test,y_train,y_test = train_test_split(x,y,test_size=0.1,random_state=2)
```

```
In [20]: print(x.shape, x_train.shape, x_test.shape)
```

```
(480, 11) (432, 11) (48, 11)
```

Training the model:

Support Vector Machine Model

```
In [21]: classifier = svm.SVC(kernel='linear')
#training the support Vector Macine model
classifier.fit(x_train,y_train)
```

```
Out[21]: SVC
SVC(kernel='linear')
```

Model Evaluation

```
In [22]: # accuracy score on training data
x_train_prediction = classifier.predict(x_train)
training_data_accaray = accuracy_score(x_train_prediction,y_train)

print('Accuracy on training data : ', training_data_accaray)
```

```
Accuracy on training data : 0.7962962962962963
```

```
In [23]: # accuracy score on training data
x_test_prediction = classifier.predict(x_test)
```

```
test_data_accuray = accuracy_score(x_test_prediction,y_test)

print('Accuracy on test data : ', test_data_accuray)
```

Accuracy on test data : 0.8125

Training Machine Learning Models

```
In [24]: #Create a fuction within many machine Learning Models
def models(x_train,y_train):

    #using Logistic Regression Algorithm to the Training set
    from sklearn.linear_model import LogisticRegression
    lr=LogisticRegression()
    lr.fit(x_train,y_train)

    #using KNeighborsClassifier Method of neighbors class to use Nearest Neighbor algo
    from sklearn.neighbors import KNeighborsClassifier
    knn=KNeighborsClassifier(n_neighbors=5,metric='minkowski',p=2)
    knn.fit(x_train,y_train)

    #using SVC method of svm class use Support Vector Machine Algorithm
    from sklearn.svm import SVC
    svc=SVC(kernel='linear',random_state=0)
    svc.fit(x_train,y_train)

    #using GaussinNB method of navie_bayes class to use Naive Bayes Algorithm
    from sklearn.naive_bayes import GaussianNB
    gnb=GaussianNB()
    gnb.fit(x_train,y_train)

    #using DecisionTreeClassifier of tree class to use Decision Tree Classifier algo
    from sklearn.tree import DecisionTreeClassifier
    dtc=DecisionTreeClassifier()
    dtc.fit(x_train,y_train)

    #using RandomForestClassifier method of ensemble class to use Random Forest Classi
    from sklearn.ensemble import RandomForestClassifier
    rfc=RandomForestClassifier(n_estimators=11,criterion='entropy',random_state=0)
    rfc.fit(x_train,y_train)

    #print model accuracy on the training data.
    print('[0] Logistic Regression Training Accuracy : ',lr.score(x_train,y_train))
    print('[1] K Nearest Neighbor Training Accuracy : ',knn.score(x_train,y_train))
    print('[2] Support Vector Machine(Linear Classifier) Training Accuracy : ',svc.score(x_train,y_train))
    print('[3] Gussian Naive Bayes Training Accuracy : ',gnb.score(x_train,y_train))
    print('[4] Decision Tree Classifier Training Accuracy : ',dtc.score(x_train,y_train))
    print('[5] Random Forest Classifier Training Accuracy : ',rfc.score(x_train,y_train))

    return lr,knn,svc,gnb,dtc,rfc
```

Evaluating Performance on Training Sets

```
In [25]: #Get and train all of the models
model=models(x_train,y_train)
```

```
[0] Logistic Regression Training Accuracy : 0.7986111111111112
[1] K Nearest Neighbor Training Accuracy : 0.7384259259259259
[2] Support Vector Machine(Linear Classifier) Training Accuracy : 0.7962962962962963
[3] Gaussian Naive Bayes Training Accuracy : 0.7916666666666666
[4] Decision Tree Classifier Training Accuracy : 1.0
[5] Random Forest Classifier Training Accuracy : 0.9837962962962963
```

Evaluating Performance on Testing Sets

```
In [26]: from sklearn.metrics import confusion_matrix
for i in range(len(model)):
    cm=confusion_matrix(y_test, model[i].predict(x_test))
    #extracting TN,FP,FN,TP
    TN,FP,FN,TP = confusion_matrix(y_test, model[i].predict(x_test)).ravel()
    print(cm)
    print('model[{}] Testing Accuracy="{ } !"'.format(i,(TP + TN) / (TP + TN + FN + FP))
    print() #Print a new line

[[ 7  8]
 [ 0 33]]
model[0] Testing Accuracy="0.8333333333333334 !"

[[ 5 10]
 [ 4 29]]
model[1] Testing Accuracy="0.7083333333333334 !"

[[ 6  9]
 [ 0 33]]
model[2] Testing Accuracy="0.8125 !"

[[ 7  8]
 [ 1 32]]
model[3] Testing Accuracy="0.8125 !"

[[ 8  7]
 [ 5 28]]
model[4] Testing Accuracy="0.75 !"

[[ 8  7]
 [ 3 30]]
model[5] Testing Accuracy="0.7916666666666666 !"
```

Saving the model

```
In [27]: import pickle
filename='Red Wine Quality Prediction.pkl'
pickle.dump(classifier, open(filename,'wb'))
```

```
In [29]: a=np.array(y_test)
predicted=np.array(classifier.predict(x_test))
loan_df.com=pd.DataFrame({"original":a,"predicted":predicted},index=range(len(a)))
loan_df.com
```


Out[29]:

	original	predicted
0	1	1
1	0	0
2	1	1
3	1	1
4	1	1
5	0	1
6	0	0
7	1	1
8	1	1
9	0	1
10	1	1
11	1	1
12	1	1
13	1	1
14	0	0
15	1	1
16	1	1
17	1	1
18	1	1
19	0	0
20	1	1
21	0	0
22	1	1
23	1	1
24	1	1
25	1	1
26	0	1
27	1	1
28	1	1
29	0	1
30	1	1
31	1	1
32	0	1

	original	predicted
33	1	1
34	1	1
35	0	1
36	1	1
37	1	1
38	0	1
39	1	1
40	1	1
41	0	1
42	1	1
43	1	1
44	0	1
45	1	1
46	1	1
47	0	0

In []: