

Importing the Required Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

Load the Dataset

```
In [2]: df=pd.read_csv('Salary_Prediction_project.csv')
df.head()
```

```
Out[2]:
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	Prof	B	19	18	Male	139750
1	Prof	B	20	16	Male	173200
2	AsstProf	B	4	3	Male	79750
3	Prof	B	45	39	Male	115000
4	Prof	B	40	41	Male	141500

```
In [3]: #check the row and column
df.shape
```

```
Out[3]: (397, 6)
```

```
In [4]: #check the information the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 397 entries, 0 to 396
Data columns (total 6 columns):
 #   Column                Non-Null Count  Dtype
---  ---
 0   rank                  397 non-null    object
 1   discipline            397 non-null    object
 2   yrs.since.phd         397 non-null    int64
 3   yrs.service           397 non-null    int64
 4   sex                   397 non-null    object
 5   salary                397 non-null    int64
dtypes: int64(3), object(3)
memory usage: 18.7+ KB
```

```
In [5]: #check the missing values
df.isnull().sum()
```

```
Out[5]: rank                0
discipline                0
yrs.since.phd            0
yrs.service              0
sex                      0
salary                  0
dtype: int64
```

```
In [6]: #check the columns  
df.dtypes
```

```
Out[6]: rank           object  
discipline          object  
yrs.since.phd       int64  
yrs.service         int64  
sex                 object  
salary              int64  
dtype: object
```

```
In [7]: df.describe()
```


```
Out[7]:
```

	yrs.since.phd	yrs.service	salary
count	397.000000	397.000000	397.000000
mean	22.314861	17.614610	113706.458438
std	12.887003	13.006024	30289.038695
min	1.000000	0.000000	57800.000000
25%	12.000000	7.000000	91000.000000
50%	21.000000	16.000000	107300.000000
75%	32.000000	27.000000	134185.000000
max	56.000000	60.000000	231545.000000

Data visualization


```
In [8]: df.salary.hist()
```

```
Out[8]: <AxesSubplot:>
```




```
In [9]: df.sex.hist()
```

```
Out[9]: <AxesSubplot:>
```



```
In [10]: sns.countplot(x='rank',data=df)
print(df['rank'].value_counts())
```

```
Prof      266
AsstProf   67
AssocProf  64
Name: rank, dtype: int64
```



```
In [11]: df.isnull().sum()
```

```
Out[11]: rank      0
discipline  0
yrs.since.phd  0
yrs.service  0
sex         0
salary      0
dtype: int64
```

```
In [12]: df.shape
```

```
Out[12]: (397, 6)
```

```
In [13]: df.dropna(how='any').shape
```

```
Out[13]: (397, 6)
```

```
In [14]: df.isnull().sum()
```

```
Out[14]: rank          0
discipline         0
yrs.since.phd      0
yrs.service         0
sex                0
salary             0
dtype: int64
```

```
In [15]: df["rank"].unique()
```

```
Out[15]: array(['Prof', 'AsstProf', 'AssocProf'], dtype=object)
```

```
In [16]: df["rank"].unique().size
```

```
Out[16]: 3
```

Handling Categorical Features or Columns

```
In [17]: df.head()
```

```
Out[17]:
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	Prof	B	19	18	Male	139750
1	Prof	B	20	16	Male	173200
2	AsstProf	B	4	3	Male	79750
3	Prof	B	45	39	Male	115000
4	Prof	B	40	41	Male	141500

```
In [18]: #Handling categorical feature ---> Gender
df['sex']=df['sex'].map({'Female':0, 'Male':1})
df.head()
```

```
Out[18]:
```

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	Prof	B	19	18	1	139750
1	Prof	B	20	16	1	173200
2	AsstProf	B	4	3	1	79750
3	Prof	B	45	39	1	115000
4	Prof	B	40	41	1	141500

```
In [19]: df['rank'].unique()
```

```
Out[19]: array(['Prof', 'AsstProf', 'AssocProf'], dtype=object)
```

```
In [20]: #Handling categorical feature ---> rank
df['rank']=df['rank'].map({'Prof':0, 'AsstProf':1, 'AssoProf':2})
df.head()
```

Out[20]:

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	0.0	B	19	18	1	139750
1	0.0	B	20	16	1	173200
2	1.0	B	4	3	1	79750
3	0.0	B	45	39	1	115000
4	0.0	B	40	41	1	141500

In [21]: `df['discipline'].unique()`

Out[21]: `array(['B', 'A'], dtype=object)`

In [22]: `df['discipline']=df['discipline'].map({'B':0,'A':1})`
`df.head()`


Out[22]:

	rank	discipline	yrs.since.phd	yrs.service	sex	salary
0	0.0	0	19	18	1	139750
1	0.0	0	20	16	1	173200
2	1.0	0	4	3	1	79750
3	0.0	0	45	39	1	115000
4	0.0	0	40	41	1	141500

Correlation Heat Map

In [23]: `plt.figure(figsize=(10,6))`
`sns.heatmap(df.corr(),cbar=True,square=True,fmt='.2f',annot=True,annot_kws={'size':8})`

Out[23]: `<AxesSubplot:>`



In []:

```
In [24]: X = df.iloc[:, :-1].values
y = df.iloc[:, 1:].values
```

```
In [25]: X.shape
y.shape
```

Out[25]: (397, 5)

```
In [32]: # Split data into training and testing
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=
```

```
In [33]: print(X.shape)
print(y.shape)
```

(442, 10)
(442,)

```
In [34]: from sklearn.ensemble import HistGradientBoostingRegressor
from sklearn.datasets import load_diabetes
X, y = load_diabetes(return_X_y=True)
est = HistGradientBoostingRegressor().fit(X, y)
est.score(X, y)
```

Out[34]: 0.9299589575098558